

DALAS Project

Drug Repurposing

Vladislav Antipin (21242244)

Paul Béglin (21408798)

Master of Computer Science, Sorbonne University

January 8, 2026

Contents

1	Introduction	2
1.1	Problem Statement	2
1.2	Motivation and Objectives	2
2	Data	2
2.1	Data Sources	2
2.2	Target variable definition	3
2.3	Feature Engineering and Description	4
2.4	Exploratory Data Analysis	5
2.4.1	Missing Values	6
2.4.2	Univariate and bivariate analysis	6
2.4.3	Dimension Reduction and Outlier Analysis	7
2.4.4	Supervised Dimension Reduction	7
2.5	Graph Representation	8
3	Methods	8
3.1	Model Overview	8
3.2	Implementation Details	9
3.3	Evaluation Strategy	9
4	Results	10
4.1	Model Comparison	10
4.2	Feature Analysis	10
4.3	Failure and Success Cases	11
5	Discussion	11
6	Conclusion	11

1 Introduction

1.1 Problem Statement

Clearly define the problem your project addresses. Explain its relevance in the scientific or practical context.

The discovery of new drugs is, although crucial, an extremely expensive and time-consuming process. Many candidate drugs ultimately turn out to be not only ineffective, but also potentially toxic. It is thus of interest to repurpose already approved drugs for diseases they have not yet been tested against. To prioritize potential clinical trials, in addition to experts' knowledge, it's important to be able to predict the likelihood of success for such studies. This project aims to explore the data and develop a potential prediction pipeline for this purpose. We acknowledge that industrial and academic drug development employ far more comprehensive methods, this project, by contrast, focuses on a preliminary, data-driven approach and is exploratory in nature.

1.2 Motivation and Objectives

Explain why this problem is important. State your main objectives and hypotheses.

The ability to predict the success of drug repurposing has academic and clinical value, it helps reducing the experimental cost and accelerating the discovery of effective treatments. The main objectives of this project are to:

- Explore public datasets to identify some relevant features
- Develop a preliminary machine learning pipeline for success prediction
- Evaluate the model performance
- Identify biases and potential ways for improvement

We decided to specify a group of diverse but closely related diseases - immune system disorders and drugs that are approved for at least one member of this group.

2 Data

2.1 Data Sources

Describe where the data comes from, collection methods, and aggregation process.

Among the vast amount of medical, pharmacological and chemical publicly available data, we focused our attention on some being able to provide us information on basic classification of diseases (MeSH), drug chemical descriptors and targets (ChEMBL), disease targets (Open Targets) as well as information on drug-disease indication studies (ClinicalTrials). Table 1 summarizes information on data sources. Although initially we were using the web scraping to get information on disease classification, we found out that this way of getting medical information is rather deprecated, so we queried the corresponding databases instead. Data concerning drugs and diseases was agglomerated using the cross references furnished by the same databases (ChEMBL and Open Targets). We used as well a UniProt identifier mapper that allowed us to unify synonymous target IDs

as well as enrich them with corresponding molecular pathway identifiers from Reactome database.

Table 1: Description of data sources.

Source	Access Method	Data Retrieved
MeSH RDF	SPARQL endpoint	Disease IDs and ontology
ChEMBL	REST API (Python client with query-style filters)	Drug chemical properties, targets and indications
UniProt	REST API (Python client)	Map target and pathway IDs
Open Targets	GraphQL API	Disease targets
ClinicalTrials.gov	REST API	Trial results and metadata

2.2 Target variable definition

Start with explaining how you got the labels.

The prediction target is a binary success / fail variable indicating whether there was shown an effect or the absence thereof of a drug against a disease. So we needed to obtain binary success labels for each clinically tried drug-disease combination. Drugs that have passed the fourth (last) clinical trial phase for a given disease are automatically labelled as successful. However, other cases are not as obvious. If the last phase wasn't (yet) passed, we can understand the reason by addressing the corresponding clinical trials. There are several possibilities: either this combination was never tried (which is most often the case); or it was tried and shown to be not-effective or effective but for earlier phases; or finally simply because there are no results published or the study was terminated for various (often financial) reasons. There exists some manually and automatically annotated clinical trial databases, but the problem is that either they are pretty limited in time span (e.g. last 10 years) or their "non-success" label encompasses simultaneously trials failed due to drug inefficiency and to other reasons. So we decided to automatically label drug-disease indications that have been tried at least once using following rule:

- If phase 4 is passed, it's definitely a success
- If there is clinical trial evidence, check out results
 - access success of the trial based on results p-values
 - trial is successful if median p-value is below 0.05, fail if above 0.2, not known otherwise
 - if the highest tried phase is at least 3 and number of successes there is twice as big as fails, assign success to the indication
 - otherwise if the latests phases has twice more fails than successes, assign fail
- If still not decided, check whether this indication was abandoned - the latest date available for any clinical trial concerning this indication was before 2020. We label such indication as fail.

It is interesting to note that most of undecided indications and clinical trials without published results happen to be approximately before 2000 and after 2010. One can

speculate that it is linked to the fact that ClinicalTrials.gov was launched in 2000 and it became mandatory to publish there all studies (but not full results) since 2007.

It’s important to note that some of those choices are arbitrary and this probably constitutes the most important weakness of this project.

2.3 Feature Engineering and Description

Provide a table or list explaining each feature and its meaning. Optionally, include relationships between features.

Description of extracted features can be found in Table 2. Categorical features were all one-hot encoded. Most drug features describe chemical and structural properties of those drugs, as well as some administration information. Among engineered features are Morgan molecular fingerprints with 2 atoms radius obtained with RDKit module from drug’s SMILES (flat string representation of the structure) and yielding a 2048 dimensional bit vector its dimension coding for presence or absence of a particular substructure. To reduce the dimensionality, first 100 dimensions were taken and called `FP_#`.

Disease features are comprised of disease ontology description, one of them being ontological category in addition to immune system diseases, called `cat_#`, for example some of immune diseases are as well neoplastic (e.g. leukemia).

For both diseases and drugs, we obtained a list of molecular (we chose protein) targets - molecules upon which a drug acts or that are involved in the pathogenesis of a disease. Since targets themselves are extremely heterogeneous and very intertwined into molecular pathways, we converted a list of targets for each drug or disease into a list of corresponding pathways. We further performed TF-IDF vectorizations on pathway lists separately for drugs and for diseases. TF-IDF helps downweighting housekeeping pathways (very common ones) and highlights informative ones. We then performed a PCA and extracted first 50 principal components for drugs and 50 for diseases, called as `drug_path_#` and `disease_path_#`. This information captures drug-drug and disease-disease mechanistic similarity based on shared pathways.

To capture the drug-disease pathway similarity, we computed a Jaccard similarity index based on corresponding pathway ID lists. We wanted to include as well drug-disease name similarity based on PubMedBERT embeddings, which reflects how related a drug and a disease in a scientific literature, but later we realized that including this feature for model training would be a mistake since PubMedBERT is trained on recent literature and thus poses a risk of a data leak.

Table 2: Description of features.

Feature	Type	Description
Drug		
drug_id	str	Unique ChEMBL identifier
chemical_probe	bool	Well-characterized and selective, typically used in biological research
dosed_ingredient	bool	Active ingredient that has been dosed in humans
inorganic_flag	bool	Inorganic molecule
natural_product	bool	Derived from a natural source (not fully synthetic)
oral	bool	Known to be administered orally
orphan	bool	Intended for use against a rare condition
parenteral	bool	Known to be administered parenterally
prodrug	bool	Prodrug, activated after administration
therapeutic_flag	bool	Has a therapeutic application, opposed to e.g. imaging
topical	bool	Known to be administered topically
veterinary	bool	Has veterinary application as well
chirality	cat	achiral,single enantiomer, mixture,nan
alogp	float	Octanol-water partition coefficient, lipophilicity indicator
aromatic_rings	int	Number of aromatic rings
full_mwt	float	Molecular weight of the full compound including any salts
hba	float	Number of hydrogen bond acceptors
hbd	float	Number of hydrogen bond donors
heavy_atoms	int	Number of heavy (non-hydrogen) atoms
mw_freebase	float	Molecular weight of parent compound
np_likeness_score	float	Natural product likeness score [Ertl et al., 2008]
num_ro5_violations	int	Violations of Lipinski’s rule-of-five, using HBA and HBD definitions
psa	float	Polar surface area
qed_weighted	float	Weighted quantitative estimate of drug likeness [Bickerton et al., 2012]
ro3_pass	int	Passes the rule-of-three (mw < 300, logP < 3 etc)
rtb	int	Number of rotatable bonds
FP_#	float	First PCs in Morgan molecular fingerprint space
drug_path_#	float	First PCs in TF-IDF space derived from pathway lists
Disease		
disease_id	str	Unique EFO identifier
ontology_depth	int	Depth in disease classification tree
nb_descendants	int	Number of descendants in disease classifications tree
nb_indications	int	Number of drug indications
cat_#	bool	Ontological category other than immune system disorders
disease_path_#	float	First PCs in TF-IDF space derived from pathway lists
Drug–Disease		
success	bool	Consensus success or failure
path_similarity	float	Jaccard similarity of pathway lists
name_similarity	float	Cosine similarity of name embeddings in PubMedBERT (later excluded due to data leak)

2.4 Exploratory Data Analysis

Discuss missing values, outliers, normalization, feature encoding, or transformations.

2.4.1 Missing Values

Among collected data and chosen features, only a particular group of drug features present missing values. To find out what it is linked to, we created an indicator variable for missingness for this group of features and built a correlation heatmap to visualize missingness patterns. Indeed, this group of chemical descriptors (such as number of aromatic rings or lipophylicity) tend to be absent together and they are associated with biotherapeutics marker. Biotherapeutics are drugs derived from living organisms rather than being chemically synthesized. Most of them are antibodies or other proteins - large biopolymers for which this group of chemical descriptors is not relevant. Indeed, manual inspection of drugs with missing values confirms this (e.g. with insulin among such drugs). So we decided to consider only small molecule drugs since chemical description of biopolymers requires another approaches, imputing those missing chemical descriptors would be a mistake. After the removal of those cases, data is ready for merging, the result containing all information on each drug-disease indication.

Table 3: Summary of datasets and merged data.

Dataset	Rows	Columns
Drugs	525	181
Diseases	149	79
Trials	14316	12
Indications	2482	12
Name embeddings	489939	5
Molecular fingerprints	1029	101
Merged data	890	252
Merged data success counts	False: 637, True: 253	

2.4.2 Univariate and bivariate analysis

Inspection of feature distributions across label values of variables not related to molecular pathways or fingerprints shows approximately bell-shaped but sometimes multimodal or slightly skewed distributions. Nevertheless, we don’t find log transformation of the features necessary here. As for disease ontologies or some drug descriptors, there are some outliers, but it may be a part of normal biological or chemical variability so we will analyze potential outliers with dimension reduction techniques. We can also notice that feature distributions are rather similar between success and failure groups. It’s expected since unfortunately such basic descriptors alone cannot discriminate success cases. However, we can notice that engineered name and path similarity features tend to be higher for success groups, this indicates that those variables might be promising predictors combined with other variables, although name similarity was excluded from further analysis since was considered as spoiler.

On a correlation heatmap for the whole merged dataset we can notice that some chemical descriptors with molecular fingerprints and drug pathway embeddings, disease ontology types with disease pathway embeddings form, as expected, clusters with strong or moderate associations. When we focus on more interpretable features (not structural or mechanistic embeddings), we can explain some of them. For example number of heavy atoms is associated with molecular weight. Number hydrogen bond acceptors and donors

(HBA and HBD) are supposed to be associated with higher hydrophilicity of a molecule and it indeed correlates positively with polar surface area (PSA) and negatively with octanol-water partition coefficient (AlogP). Number of rule of 5 violation is a feature based on some of those chemical descriptors so it is, as expected, associated with them. Interestingly, number of aromatic rings correlates negatively with the indicator that the drug was obtained from the natural source (NP-likeness score). Apparently, purely synthetic drugs contain more aromatic rings.

As for disease ontology, we can for example see that immune diseases labelled as hemic or lymphatic tend to be neoplastic as well, it makes sense since most of oncological conditions concerning the immune system are indeed various types of leukemias and lymphomas. Another example is positive association between urogenital diseases and infections.

2.4.3 Dimension Reduction and Outlier Analysis

When we perform PCA on centered and scaled merged data we can first notice that first two principal components explain only 10 % of total variance, and most of the variance is distributed among a large number of components. It can mean that no linear combination of these features form a single or several dominant axes of data variability. This suggest that the feature space is intrinsically high dimensional and associations between variables are likely non-linear which is expected for such a complex system. Furthermore, a lot of features were explicitly designed to be uncorrelated (principal components for molecular fingerprints and pathways embeddings). We can however cautiously take a look at the data projection on first two principal components and notice that the major variation is associated with disease heterogeneity, since the most prominent arrows on correlation plot associated with first principal components are disease category indicators and disease pathway embeddings, for example neoplastic blood disorders as opposed to respiratory tract diseases.

The second principal component is rather associated with drug related variability, in particular we can notice a group of potential outliers for high values of first and second principal component. Their closer inspection reveals that these indications involve three drugs sirolimus, everolimus and ridaforolimus. These drugs turn out to be antineoplastic and immunosuppressor drugs with the same structural code - large macrolide ring. We considered it as a part of normal drug variation that cannot be attributed to a mistake or poor data filtering and thus we decided to not exclude them.

Although PCA is unsupervised method, when we label the data a posteriori, we can see a slight separation of success and fail cases along the first principal component. We could speculate that it might be linked to different success rates for various disease types, for example clinical studies for neoplastic blood diseases might have fewer successes because of smaller success rate or may be because they are more often studied...

Non linear dimension reduction methods such as UMAP and t-SNE result in a rather large number of small clusters. Apparently, those clusters correspond to the indications involving same or very similar drugs or diseases.

2.4.4 Supervised Dimension Reduction

By performing supervised dimension reduction models - LDA and PLS-DA we can see that the label distributions along LDA axis or in projection to first two components of PLS-DA, are, although overlapping, are rather well separated

2.5 Graph Representation

We need to acknowledge that the chosen approach for binary classification of drug-disease indications, although is valid, doesn’t model the true structure of the problem.

Indeed, the data is inherently structured and a bipartite drug-disease graph with a connecting layer comprised of associated molecular pathways. And most of advanced research in the field of drug discovery involve working with graphs and Graph Neural Networks directly rather than engineering features as we did. To this end, we organized the same data in the form of `HeteroData` graph with Pytorch Geometric module. Although we think that the richer information about connections on a molecular / mechanistic levels as well as more observations are needed to build a graph suitable for the robust model training.

```
HeteroData(  
    drug={ x=[525, 127] },  
    disease={ x=[149, 25] },  
    pathway={ x=[1487, 1] },  
    (drug, interacts, pathway)={ edge_index=[2, 17338] },  
    (disease, associated_with, pathway)={ edge_index=[2, 12064] },  
    (drug, tried_against, disease)={  
        edge_index=[2, 890],  
        edge_label=[890],  
    }  
)
```

3 Methods

3.1 Model Overview

Briefly describe the models used. Include general principles (e.g., linear models, random forests, neural networks).

Since the prediction target is a binary success/failure for given drug-disease indication, we evaluated several supervised models for binary classification, including Logistic Regression, k-Nearest Neighbors (kNN), Support Vector Machines (SVM), Random Forest, Gradient Boosting, and XGBoost.

Logistic regression estimates probability of success using sigmoid-transformed linear combinations of the input features. Its biggest advantage is high interoperability since feature weights indicate their importance, however it is limited to linearly separable problems.

k-Nearest Neighbors is instance based model, making predictions by looking at labels of the closest points in the feature space. Therefore, there is no actual training phase and the assumption behind this model is that similar observations tend to have the same label. Nevertheless it is very sensitive to noise and the curse of dimensionality, as well as poorly interpretable (only by samples that voted).

Support Vector Machines are fitting a decision boundary, maximizing the margin between classes. Although it may be similar to regularized linear regression, the usage of kernel functions allows fitting of non-linear decision boundaries.

Tree-based ensemble learning models leverage hierarchical decision rules and combine prediction of multiple models. A random forest aggregates multiple decision trees based on random subsets of variables and observations, this decreases the overfitting so characteristic of decision trees. Gradient boosting and XGBoost instead build trees sequentially, where each new tree learns errors of the previous ones. XGboost is an optimized version of Gradient Boosting, but has more parameters to tune.

As for graph-based approach, we chose to use a Heterogenous Graph Neural Network (GNN) link classifier from PyTorch Geometric module. Instead of learning patterns based on tabular features only, it learns representations by propagating information across connected nodes and edges, enabling the model to capture more complex relationships. This approach avoids manual feature engineering, but requires a large number of samples and rich informative connections in a graph, and, as other Artificial Neural Networks, prone to overfitting.

3.2 Implementation Details

Mention important parameters, software/libraries used.

All numeric features are centered and scaled. No imputation was needed since observations with missing values corresponding to biopolymer drugs were excluded.

Most of the mentioned models are imported from `sklearn` module, except for XGBoost imported from `xgboost` and GNN from `torch_geometric` modules. In order to perform hyperparameter tuning, we used a grid-search cross validation `GridSearchCV` with search spaces indicated in Table 4.

Since most of the clinical trials and thus corresponding indications usually fail, the classes are imbalanced with failures being several times more abundant (around 3 times for our data). So we used options for class balancing where it was possible (all models except for kNN and Gradient Boosting).

Table 4: Hyperparameter search spaces for each classifier.

Model	Hyperparameters
k-Nearest Neighbors	<code>n_neighbors</code> $\in \{5, 10, 20, 30\}$, <code>weights</code> $\in \{\text{uniform}, \text{distance}\}$, <code>p</code> $\in \{1, 2\}$
Logistic Regression	<code>C</code> $\in \{0.01, 0.1, 1, 10, 100\}$, <code>l1_ratio</code> $\in \{0, 0.25, 0.5, 0.75, 1\}$
SVM (linear)	<code>C</code> $\in \{0.01, 0.1, 1, 10\}$
SVM (polynomial)	<code>C</code> $\in \{0.1, 1, 10\}$, <code>degree</code> $\in \{2, 3, 4\}$, <code>gamma</code> $\in \{\text{scale}, \text{auto}\}$
SVM (RBF)	<code>C</code> $\in \{0.1, 1, 10\}$, <code>gamma</code> $\in \{\text{scale}, \text{auto}\}$
Random Forest	<code>n_estimators</code> $\in \{100, 300, 500\}$, <code>max_depth</code> $\in \{\text{None}, 10, 20\}$, <code>min_samples_split</code> $\in \{2, 5\}$, <code>min_samples_leaf</code> $\in \{1, 2\}$
Gradient Boosting	<code>n_estimators</code> $\in \{100, 300\}$, <code>max_depth</code> $\in \{3, 5\}$, <code>learning_rate</code> $\in \{0.01, 0.05, 0.1\}$, <code>subsample</code> $\in \{0.8, 1.0\}$
XGBoost	<code>n_estimators</code> $\in \{100, 300\}$, <code>max_depth</code> $\in \{3, 5\}$, <code>learning_rate</code> $\in \{0.01, 0.05, 0.1\}$, <code>subsample</code> $\in \{0.8, 1.0\}$, <code>colsample_bytree</code> $\in \{0.8, 1.0\}$, <code>gamma</code> $\in \{0, 0.1\}$

3.3 Evaluation Strategy

Drug-disease indication success prediction requires avoiding temporal data leak - in order to simulate fairly the real application, we need to use older research results to predict

success of new indications. Thus, we performed a train-test split based on on start date of the first clinical trial for each indication. Taking 15% for test set results in temporal splitting by August 2014.

The training set is further used for hyperparameter tuning using grid search with 5-fold cross-validation and model selection based on cross-validation performances.

Since we deal with imbalanced classes, we must not use accuracy as evaluation metric since it will underemphasize incorrect prediction of minority (success) class. We thus used F1 score (harmonic mean of precision and recall) for model hyperparameter tuning, selection and evaluation, as well as precision-recall curves instead of ROC curves for analyzing how clarification result varies based on decision threshold (it works only for models yielding class probabilities).

4 Results

4.1 Model Comparison

Compare models quantitatively (accuracy, RMSE, AUC, etc.) and qualitatively (strengths/weaknesses).

Best estimators after hyperparameter tuning are presented in Table 5. Based on mean cross-validation scores, we can say that all models perform relatively similarly having mean F1 score between 0.6 and 0.65. The best one is XGBoost and the worst one is kNN, what is expected given high dimensionality of the problem.

Table 5: Model comparison summary.

Model	Mean F1 score	Std F1 score
XGBoost	0.656383	0.038165
Logistic Regression	0.619974	0.037465
Support Vector Machine	0.619910	0.033970
Gradient Boost	0.609287	0.092787
Random Forest	0.601969	0.044340
k-Nearest Neighbors	0.591973	0.052311

We will use then XGBoost for final evaluation on test set and we obtain precision 0.658, recall 0.595 and F1 score 0.625 (see Figure).

Precision-recall curve presents rather smooth descent from 1 to chance level (success prevalence of 0.313). It means that there are no abrupt drops in precision as model becomes less strict (higher recall) and that model performance is rather stable.

4.2 Feature Analysis

Analyze the impact of key features on model performance (e.g., feature importance, correlation).

Since XGBoost model is not interpretable by design, we need to use post-hoc explanation strategy. We chose to use a permutation feature importance estimation (still using F1 score) and partial dependence analysis from `sklearn` module. Permutation of

pathway similarity feature results in 15% decrease in F1 score, around 5% for number of indications for a disease (proxy of disease "popularity") and 6-th component in disease pathway embedding space. Other features, including various disease and drug pathway embeddings and molecular fingerprints present less than 5% decrease in F1 score.

Partial dependence plot shows that the steepest increase in success prediction happens at around 0.3 for pathway similarity and 90 for number of indications for a disease. Two dimensional partial dependence plot for pathway similarity and number of indications shows that the model needs simultaneously high pathway similarity and large number of indications to be sure in success prediction.

4.3 Failure and Success Cases

Highlight scenarios where the models perform well or poorly. Include illustrative figures or tables.

We added to the test set features indicating whether each observation is true positive (TP), true negative (TN), false positive (FP) or false negative (FN) and just correct (TP or TN).

We can analyze Pearson correlations of those confusion estimates with predictor features (being cautious with interpretation since these introduced variables are binary labels, so it must be interpreted as point-biserial correlation) as well as distributions of predictors in correctly classified and misclassified cases. We can see that there is no major difference between those cases. The distributions for some variables (among which there are number of indications, some categories of diseases, QED drug likeness) present a slight bias towards higher feature values for correctly predicted cases, but their interpretation will be speculative.

5 Discussion

Interpret the results. Relate findings to your original objectives. Discuss limitations and possible improvements.

6 Conclusion

Summarize your work, key findings, and potential future directions.

References

- [Bickerton et al., 2012] Bickerton, G. R., Paolini, G. V., Besnard, J., Muresan, S., and Hopkins, A. L. (2012). Quantifying the chemical beauty of drugs. *Nat Chem*, 4(2):90–98.
- [Ertl et al., 2008] Ertl, P., Roggo, S., and Schuffenhauer, A. (2008). Natural product-likeness score and its application for prioritization of compound libraries. *J. Chem. Inf. Model.*, 48(1):68–74. Epub 2007 Nov 23.